

Founding Engineer Technical Exercise

Proposal Ingestion → Invitation to Bid Automation

Context

Bridgeline Technologies is building AI-powered software to eliminate manual workflows in pre-construction. One of the most painful tasks for General Contractors is extracting subcontractor contact information from proposals (PDFs, spreadsheets, emails) and manually re-entering it into Invitation to Bid (ITB) systems.

This exercise simulates a real problem we solve.

Objective

Build a **simple web application** that:

1. Accepts a set of subcontractor proposals
2. Extracts key contact information
3. Automatically populates that data into an **Invitation to Bid** interface
4. Information about invitation to bid can be found here:
https://en.wikipedia.org/wiki/Invitation_to_tender

The goal is not perfection—it's to demonstrate how you think, architect, and ship.

Input

You will be provided with:

- A folder of **mock subcontractor proposals** (mixed formats)
 - PDFs
 - Excel files
 - (Optional) plain text or email-style docs

Each proposal may contain:

- Company name
- Contact name
- Email address
- Phone number
- Trade / scope

Expect inconsistency and messiness—this is intentional.

Requirements

1. Proposal Upload

- A basic UI that allows users to upload one or more proposal files
- Files should be processed server-side

2. Data Extraction

- Extract the following fields where possible:
 - Company Name
 - Contact Name
 - Email
 - Phone
 - Trade / Scope
- You may use:
 - Rule-based parsing
 - OCR
 - LLMs
 - Or any hybrid approach
- Accuracy matters, but **transparency matters more**
 - If confidence is low, show it

3. Review & Edit Step

- Display extracted data in a simple table or form
- Allow the user to:
 - Edit fields
 - Confirm or reject entries
- This mirrors real pre-construction workflows

4. Invitation to Bid Creation

- Once confirmed, populate the data into a mock **Invitation to Bid** screen:
 - Subcontractor list
 - Contact info
 - Trade assignment
 - No need to send real emails—just show the populated state
-

Technical Expectations

- Use any modern web stack you're comfortable with
 - Clean, readable, production-minded code
 - Reasonable data models
 - Thoughtful error handling
 - Clear README explaining:
 - Architecture choices
 - Tradeoffs
 - What you'd improve with more time
-

What We're Evaluating

We care less about polish and more about **how you think**:

- How you handle unstructured, real-world data
 - Product intuition (what the user needs vs. what's technically interesting)
 - Ability to make pragmatic tradeoffs
 - Code clarity and extensibility
 - How you communicate assumptions and limitations
-

Bonus (Optional)

- Confidence scoring on extracted fields
- Deduplication of subcontractors
- Handling multiple contacts per company
- Notes on how this would scale in production